# Effectively using GPUs with Julia

Tim Besard (@maleadt)





# What you should know



- High-level programming language with low-level performance
- Solves "two language problem", but requires proficiency



- Hardware accelerator for massively parallel applications
- Throughput oriented: hard to program

# Why not both?



- High-level programming without GPU experience
- Low-level programming for high-performance and flexibility

# Choice of hardware







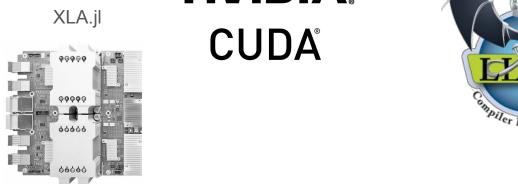


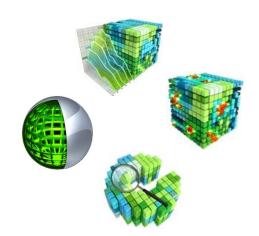




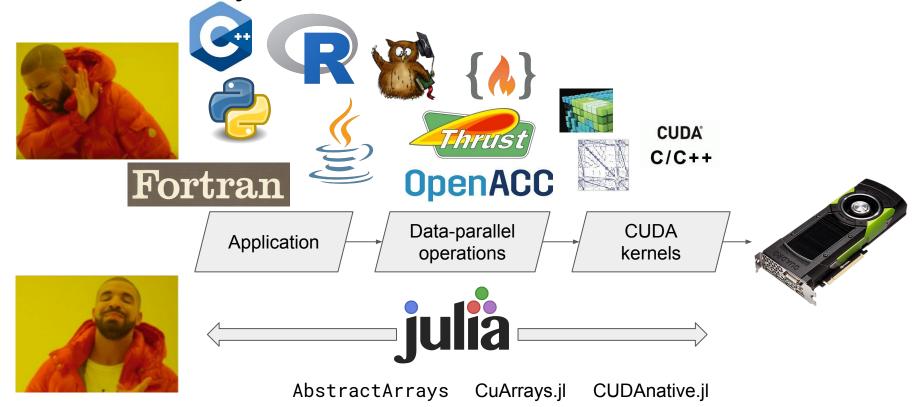









# How to train your GPU: 10.000 foot view





## Hello GPU!

# Code is specialized

# Code is compiled

```
julia> @device_code_llvm @cuda say(42)
define void @sav(i64) {
entry:
 %1 = call i32 @llvm.nvvm.read.ptx.sreg.tid.x()
%addconv = add nuw nsw i32 %1. 1
 %2 = zext i32 %addconv to i64
%3 = alloca %printf_args.0
 %4 = bitcast %printf_args.0* %3 to i8*
 %5 = getelementptr inbounds %printf_args.0,
      %printf_args.0* %3, i64 0, i32 0
 store i64 %2. i64* %5
 %6 = getelementptr inbounds %printf_args.0,
      %printf args.0* %3. i64 0. i32 1
 store i64 %0. i64* %6
%7 = call i32 @vprintf(i8* ..., i8* %4)
 ret void
```

# Code is compiled

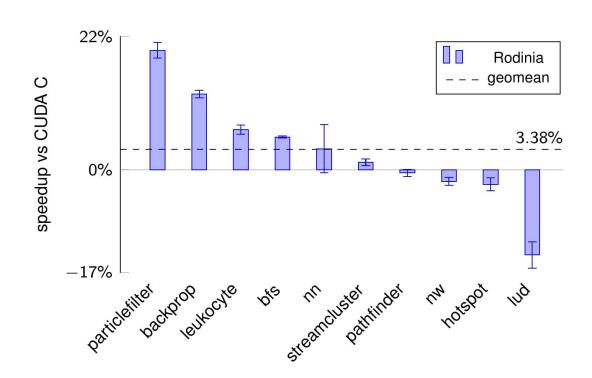
```
pkg> add CUDAnative
julia> using CUDAnative
julia> function say(num)
        @cuprintf("Thread %ld says: %ld\n",
                  threadIdx().x, num)
        return
      end
julia> @cuda threads=4 say(42)
Thread 1 says: 42
Thread 2 says: 42
Thread 3 says: 42
Thread 4 says: 42
```

```
julia> @device_code_sass @cuda say(42)
.say:
 S2R R1, SR_TID.X;
 IADD32I R1. R1. 0x1:
 MOV R2. c[0x0][0x44]:
 IADD32I R2. R2. -0x10:
 MOV R8, c[0x0][0x140];
 MOV R9, c[0x0][0x144];
 MOV R3. RZ:
 MOV R7. RZ:
 MOV32I R4, 32@lo(__unnamed_1);
 STL.64 [R2+0x8]. R8:
 LOP.OR R6, R2, c[0x0][0x24];
 MOV32I R5, 32@hi(__unnamed_1);
 STL.64 [R2]. R1:
 JCAL `(vprintf);
 MOV RZ, RZ:
 EXIT;
```

# Why should you care?

1) Performance

2) Powerful abstractions




# Show me what you got

# Show me what you got

```
julia> a = CuArray([1., 2., 3.])
julia> function apply(op, a)
        i = threadIdx().x
        a[i] = op(a[i])
        return
     end
julia> @cuda threads=length(a) Map(x->x^2, a)
julia> a
3-element CuArray{Float32,1}:
1.0
4.0
9.0
```

```
julia> @device_code_ptx @cuda apply(x->x^2, a)
apply(.param .b8 a[16])
                      %rd1, [a+8];
       ld.param.u64
       mov.u32
                      %r1, %tid.x;
       // index calculation
       mul.wide.u32
                     %rd2, %r1, 4;
       add.s64
                      %rd3, %rd1, %rd2;
                             %rd4. %rd3:
       cvta.to.global.u64
       ld.global.f32 %f1, [%rd4];
       mul.f32
               %f2, %f1, %f1;
       st.global.f32
                     [%rd4], %f2;
       ret:
```



# Not just another array library

```
julia> a = CuArray([1,2,3])
3-element CuArray{Int64,1}:
1
2
3
```



dot syntax

```
julia> function apply(op, a)
        i = threadIdx().x
        a[i] = op(a[i])
       end
julia> @cuda threads=length(a) apply(op, a)
julia> map(op, a)
julia> reduce(binop, a)
julia> broadcast(+, [1], [2 2], [3 3; 3 3])
2×2 CuArray{Int64,2}:
6 6
6 6
julia> [1] .+ [2 2] .+ [3 3; 3 3]
2×2 CuArray{Int64,2}:
6 6
```

# Not just another array library

```
julia> a = CuArray([1f0, 2f0, 3f0])
3-element CuArray{Float32,1}:
1.0
2.0
3.0

julia> f(x) = 3x^2 + 5x + 2

julia> a .= f.(2 .* a.^2 .+ 6 .* a.^3 .- sqrt.(a))
3-element CuArray{Float32,1}:
    184.0
    9213.753
96231.72
```

### Single kernel!

- Fully specialized
- Highly optimized
- Great performance

# Vendor libraries



```
julia> a = CuArray{Float32}(undef, (2,2));
                                                  CUFFT
                                                  julia> CUFFT.plan_fft(a) * a
CURAND
                                                  2-element CuArray{Complex{Float32},1}:
julia> rand!(a)
                                                   -1.99196+0.0im 0.589576+0.0im
2×2 CuArray{Float32,2}:
                                                   -2.38968+0.0im -0.969958+0.0im
0.73055 0.843176
0.939997 0.61159
                                                  CUDNN
                                                  julia> softmax(real(ans))
CUBLAS
                                                  2×2 CuArray{Float32,2}:
julia> a * a
                                                   0.15712 0.32963
2×2 CuArray{Float32,2}:
                                                  0.84288 0.67037
 1.32629 1.13166
 1.26161 1.16663
                                                  CUSPARSE
                                                  julia> sparse(a)
CUSOLVER
                                                  2×2 CuSparseMatrixCSR{Float32,Int32}
julia> LinearAlgebra.gr!(a)
                                                  with 4 stored entries:
CuQR{Float32, CuArray{Float32,2}}
                                                    [1, 1] = -1.1905
with factors Q and R:
                                                    [2. 1]
                                                           = 0.489313
Float32[-0.613648 -0.78958; -0.78958 0.613648]
                                                    [1, 2] = -1.00031
Float32[-1.1905 -1.00031; 0.0 -0.290454]
                                                           = -0.290454
```

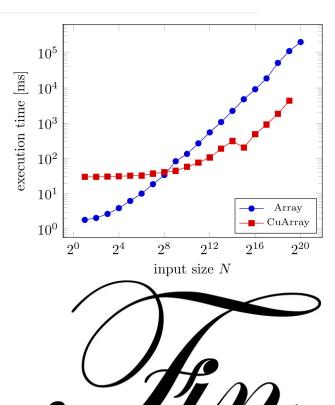
# Effective GPU Programming

How do you actually use this stuff?

### Types and gradients, including Forward.gradient

https://discourse.julialang.org/t/types-and-gradients-including-forward-gradient/946

```
using LinearAlgebra
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
loss \nabla w(w, b, x, y) = ...
lossdb(w, b, x, y) = ...
function train(w, b, x, y ; lr=.1)
   w = lmul!(lr, loss \nabla w(w, b, x, y))
   b = 1r * lossdb(w, b, x, y)
   return w. b
end
n = 100
p = 10
x = randn(n,p)'
                                                   x = CuArray(x)
y = sum(x[1:5,:]; dims=1) .+ randn(n)'*0.1
                                                   y = CuArray(y)
w = 0.0001*randn(1,p)
                                                   w = CuArray(w)
b = 0.0
for i=1:50
   w, b = train(w, b, x, y)
end
```



### cuArrays vs CUDANative

■https://discourse.julialang.org/t/cuarrays-vs-cudanative/17504

```
function diff_y(a, b)
   s = size(a)
   for j = 1:s[2]
       for i = 1:s[1]
           @inbounds a[i,j] = b[i,j+1] - b[i,j]
       end
  end
end
N = 64
nx = N^2
ny = N
a = ones(Float32, nx, ny-1)
b = ones(Float32, nx, ny)
julia> using BenchmarkTools
julia> @btime diff_y($a,$b);
 39.599 \mus (0 allocations: 0 bytes)
julia> @btime diff_y($(CuArray(a)),$(CuArray(b)));
4.499 s (3354624 allocations: 165.38 MiB)
```

### Performance killers

#### 1. Scalar iteration is slooooow

```
function diff_y(a, b)
   s = size(a)
   for j = 1:s[2]
       for i = 1:s[1]
           @inbounds a[i,j] = b[i,j+1] - b[i,j]
       end
   end
end
julia> CuArrays.allowscalar(false)
julia> diff_y(CuArray(a), CuArray(b))
ERROR: scalar getindex is disallowed
Stacktrace:
[5] getindex at ./abstractarray.jl
[6] diff_y(::CuArray, ::CuArray)
    at ./REPL[109]:5
. . .
```

```
function diff_y(a, b)
   s = size(a)
  for i = 1:s[2]
      @inbounds a[:,j] = b[:,j+1] - b[:,j]
  end
end
julia> @btime diff_y($(CuArray(a)),$(CuArray(b)));
2.503 ms (16884 allocations: 661.50 KiB)
```

21

### Performance killers

### 2. Avoid multiple kernels

## Performance killers

### 3. Bump the problem size

```
julia> N = 256
julia> @btime diff_y($(CuArray(a)),$(CuArray(b)));
1.494 ms (40 allocations: 2.08 KiB)
julia> @btime diff_y($a,$b);
11.719 ms (2 allocations: 128 bytes)
```

### 4. Keep data on the GPU

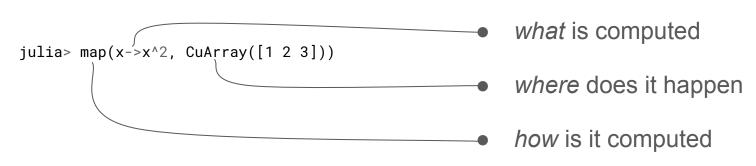
```
julia> @btime diff_y(CuArray($a),CuArray($b));
72.050 ms (93 allocations: 255.50 MiB)
```

# Strenghts

- 1. Single, productive programming language
- 2. Platform-independent, generic code
- 3. High-level, zero-cost abstractions
- 4. Great performance potential
- 5. Composability
- 6. Optimizability

# Composability

### **Separation of concerns**



CUDAnative.jl 2383 LOC GPUArrays.jl 1468 LOC CuArrays.jl 859 LOC (without libraries)

# Composability: reuse of libraries

```
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
julia> loss(w,b,x,y)
4.222961132618434
julia> loss\nabla w(w, b, x, y)
1×10 CuArray{Float64,2}:
-1.365 -1.961 -1.14 -2.023 -1.981 -0.2993 -0.2667 -0.07669 -1.038 -0.1823
using ForwardDiff
loss\nabla w(w, b, x, y) = ForwardDiff.gradient(w -> loss(w, b, x, y), w)
julia> @which mul!(w, w, x)
mul!(...) in CuArrays.CUBLAS at src/blas/highlevel.jl
julia> @which mul!(w, w, ForwardDiff.Dual.(x))
mul!(...) in CuArrays at src/generic_matmul.jl
```

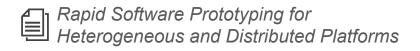
# But 18:00 Resident of the Resi

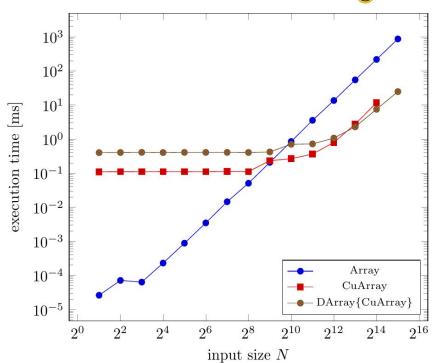
# Composability: reuse of infrastructure



```
julia> A = rand(4096,4096)
4096×4096 Array{Float64,2}
```

### JuliaParallel / DistributedArrays.jl





# Optimizability: it's julia all the way down

```
function seed!(duals::AbstractArray{Dual{T,V,N}}, x, seed::Partials{N,V} where {T,V,N}
    in eachindex(duals)
      duals[i] = Dual{T,V,N}(x[i], seed)
   return duals
end
function ForwardDiff.seed!(duals::AbstractArray{Dual{T,V,N}}, x, seed::Partials{N,V} where {T,V,N}
   duals .= Dual{T,V,N}.(x, Base.RefValue(seed))
   return duals
end
function ForwardDiff.seed!(duals::CuArray{Dual{T,V,N}}, x, seed::Partials{N,V} where {T,V,N}
   function kernel(duals, x, seed)
       i = threadIdx().x
       duals[i] = Dual\{T, V, N\}(x[i], seed)
       return
   end
   @cuda threads=length(duals) kernel(duals, x, seed)
   return duals
end
```

# Optimizability: it's julia all the way down

- Rewrite using array abstractions using CuArrays + generic code
- 2. Avoid GPU antipatterns
- 3. Specialize with broadcast expressions
- 4. Specialize with GPU kernels



### 1. Reflection and introspection

```
julia> using CUDAnative
julia> @device_code_llvm curand(2) .+ 2
define void @ptxcall_anonymous(...) {
@device_code_{lowered, typed, warntype, llvm, ptx, sass}
julia> ENV["JULIA_DEBUG"] = "CUDAnative"
julia > curand(2) .+ 2;
 Debug: Compiled getfield(GPUArrays, ...)() to PTX 3.5.0 for SM 3.5.0 using 8 registers.
 Memory usage: 0 bytes local, 0 bytes shared, 0 bytes constant
 @ CUDAnative ~/Julia/CUDAnative/src/execution.jl
```

### 2. Performance measurements

```
julia> const x = CuArray{Float32}(undef, 1024)

julia> using BenchmarkTools
julia> @benchmark CuArrays.@sync(identity.($x))
BenchmarkTools.Trial:
  memory estimate: 1.34 KiB
  allocs estimate: 33
    -----
  minimum time: 13.824 μs (0.00% GC)
  median time: 16.361 μs (0.00% GC)
  mean time: 16.489 μs (0.00% GC)
  maximum time: 401.689 μs (0.00% GC)
  ------
  samples: 10000
  evals/sample: 1
```

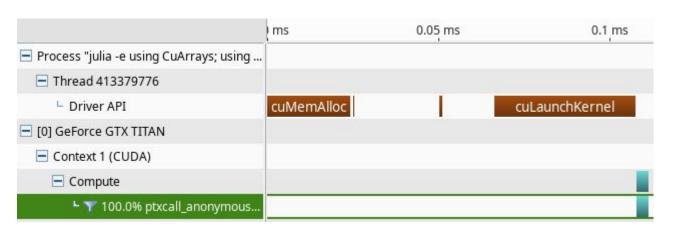
### 2. Performance measurements

### 2. Performance measurements

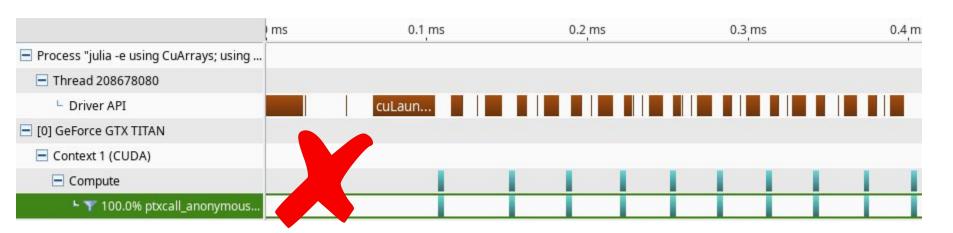
```
julia> const x = CuArray\{Float32\}(undef, 1024)
julia> using BenchmarkTools
julia> @benchmark CuArrays.@sync(identity.($x))
                                                        Accurate measurements of
BenchmarkTools.Trial:
                                                         possible short-running code
                  13.824 µs (0.00% GC)
minimum time:
                  401.689 µs (0.00% GC)
maximum time:
julia> CuArrays.@time CuArrays.@sync identity.(x);
                                                        Memory allocation behavior
0.000378 seconds (57 CPU allocations: 1.938 KiB)
                  (1 GPU allocation: 4.000 KiB)
julia> using CUDAdrv
                                                        Application performance metrics
julia> CUDAdrv.@elapsed identity.(x)
5.888f-6
```

```
$ nvprof --profile-from-start off julia
julia > const x = CuArray{Float32}(undef, 1024)
julia> identity.(x)
julia > CUDAdrv.@profile begin
        identity.(x)
      end
julia> exit()
==22272== Profiling result:
           Type Time(%)
                                       Calls
                                                             Min
                                                                      Max Name
                              Time
                                                   Avg
GPU activities: 100.00% 3.5520us
                                             3.5520us 3.5520us 3.5520us
                                                                          ptxcall_anonymous
    API calls:
                 61.70% 39.212us
                                             39.212us 39.212us 39.212us cuLaunchKernel
                 37.36% 23.745us
                                             23.745us
                                                       23.745us 23.745us cuMemAlloc
                  0.93%
                            592ns
                                                296ns
                                                          222ns
                                                                   370ns cuCtxGetCurrent
```

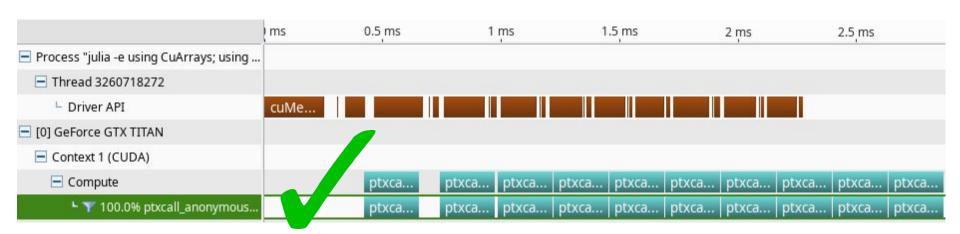
```
$ nvvp julia
julia> identity.(CuArray{Float32}(undef, 1024))
```

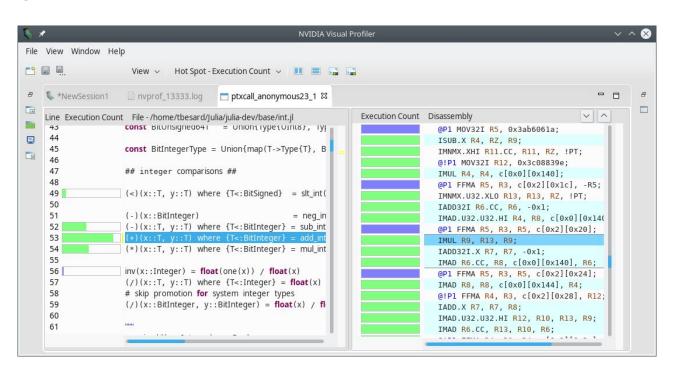


```
$ nvvp julia
julia> identity.(CuArray{Float32}(undef, 1024))
```



```
$ nvvp julia
julia> sin.(CuArray{Float32}(undef, 1024, 1024))
```





# Conclusion

- Great tools for single-language GPU programming
  - CuArrays.jl: high-level and productive
  - CUDAnative.jl: low-level performance

- Strengths: optimizability & composability
- Weaknesses: run into GPU limitations

- Tools
- Community: Slack and Discourse

# Effectively using GPUs with Julia

Tim Besard (@maleadt)

http://julialang.slack.com/ https://discourse.julialang.org/c/domain/gpu



